

# Convolutional Neural Networks Project: Street View Housing Number Digit Recognition

## Objective

To build a CNN model that can recognize the digits in the images.

## Dataset

Here, I will use a subset of the original data to save some computation time. The dataset is provided as a .h5 file. The basic preprocessing steps have been applied on the dataset.

## Mount the drive

I started by mounting the Google drive.

In [26]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## Importing the necessary libraries

In [27]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

import tensorflow as tf

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, BatchNormalization, Dropout, Flatten, LeakyReLU
from tensorflow.keras.losses import categorical_crossentropy
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
```

Let us check the version of tensorflow.

In [28]:

```
print(tf.__version__)

2.8.0
```

## Load the dataset

- load the dataset that is available as a .h5 file.
- Split the data into the train and the test dataset.

In [30]:

```
import h5py

# Open the file as read only
# can make changes in the path as required

h5f = h5py.File('/content/SVHN_single_grey1.h5', 'r')
```

*# Load the the train and the test dataset*

```
X_train = h5f['X_train'][:]
y_train = h5f['y_train'][:]
X_test = h5f['X_test'][:]
y_test = h5f['y_test'][:]
```

*# Close this file*

```
h5f.close()
```

check the number of images in the training and the testing dataset.

```
len(X_train), len(X_test)
```

(42000, 18000)

**Observation:**

- There are 42,000 images in the training data and 18,000 images in the testing data.

## Visualizing images

- Use X\_train to visualize the first 10 images.
- Use Y\_train to print the first 10 labels.

visualizing the first 10 images in the dataset

*# Visualizing the first 10 images in the dataset and printing their labels*

**%matplotlib inline**

**import matplotlib.pyplot as plt**

plt.figure(figsize = (10, 1))

**for i in range(10):**


    plt.subplot(1, 10, i+1)

    plt.imshow(X\_train[i], cmap = "gray") *# Write the function to visualize images*

    plt.axis('off')

plt.show()

print('label for each of the above image: %s' % (y\_train[0:10]))



label for each of the above image: [2 6 7 4 4 0 3 0 7 3]

## For Data preparation

- Print the shape and the array of pixels for the first image in the training dataset.
- Reshape the train and the test dataset because we always have to give a 4D array as input to CNNs.
- Normalize the train and the test dataset by dividing by 255.
- Print the new shapes of the train and the test dataset.
- One-hot encode the target variable.

*# Shape and the array of pixels for the first image*

print("Shape:", X\_train[0].shape)

print()

print("First image:\n", X\_train[0])

Shape: (32, 32)

First image:

```
[[ 33.0704 30.2601 26.852 ... 71.4471 58.2204 42.9939]
 [ 25.2283 25.5533 29.9765 ... 113.0209 103.3639 84.2949]
 [ 26.2775 22.6137 40.4763 ... 113.3028 121.775 115.4228]
 ...
 [ 28.5502 36.212 45.0801 ... 24.1359 25.0927 26.0603]
 [ 38.4352 26.4733 23.2717 ... 28.1094 29.4683 30.0661]
 [ 50.2984 26.0773 24.0389 ... 49.6682 50.853 53.0377]]
```

In [34]:

```
# Reshaping the dataset to be able to pass them to CNNs. Remember that we always have to give a 4D array as input to CNNs
```

```
X_train = X_train.reshape(X_train.shape[0], 32, 32, 1)
```

```
X_test = X_test.reshape(X_test.shape[0], 32, 32, 1)
```

In [35]:

```
# Normalize inputs from 0-255 to 0-1
```

```
X_train = X_train / 255.0
```

```
X_test = X_test / 255.0
```

In [36]:

```
# New shape
```

```
print('Training set:', X_train.shape, y_train.shape)
```

```
print('Test set:', X_test.shape, y_test.shape)
```

```
Training set: (42000, 32, 32, 1) (42000,)
```

```
Test set: (18000, 32, 32, 1) (18000,)
```

**encode the labels in the target variable y\_train and y\_test**

In [37]:

```
# Write the function and appropriate variable name to one-hot encode the output
```

```
y_train = tf.keras.utils.to_categorical(y_train)
```

```
y_test = tf.keras.utils.to_categorical(y_test)
```

```
# test labels
```

```
y_test
```

Out[37]:

```
array([[0., 1., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 1., 0., 0.],
       [0., 0., 1., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 1., 0., 0.],
       [0., 0., 0., ..., 0., 0., 1.],
       [0., 0., 1., ..., 0., 0., 0.]], dtype=float32)
```

**Observation:**

- Notice that each entry of the target variable is a one-hot encoded vector instead of a single label.

## Model Building

Now that I have done data preprocessing, let's build a CNN model.

In [38]:

```
# Fixing the seed for random number generators
```

```
np.random.seed(42)
```

```
import random
```

```
random.seed(42)
```

```
tf.random.set_seed(42)
```

Model Architecture

- **function** that returns a sequential model with the following architecture:
  - First Convolutional layer with **16 filters and the kernel size of 3x3**. Use the **'same' padding** and provide the **input shape = (32, 32, 1)**
  - Add a **LeakyRelu layer** with the **slope equal to 0.1**
  - Second Convolutional layer with **32 filters and the kernel size of 3x3 with 'same' padding**
  - Another **LeakyRelu** with the **slope equal to 0.1**
  - A **max-pooling layer** with a **pool size of 2x2**
  - **Flatten** the output from the previous layer
  - Add a **dense layer with 32 nodes**
  - Add a **LeakyRelu layer with the slope equal to 0.1**
  - Add the final **output layer with nodes equal to the number of classes, i.e., 10** and **'softmax' as the activation function**
  - Compile the model with the **loss equal to categorical\_crossentropy, optimizer equal to Adam(learning\_rate = 0.001), and metric equal to 'accuracy'**. Do not fit the model here, just return the compiled model.
- Call the function `cnn_model_1` and store the output in a new variable.
- Print the summary of the model.
- Fit the model on the training data with a **validation split of 0.2, batch size = 32, verbose = 1, and epochs = 20**. Store the model building history to use later for visualization.

Build and train a CNN model as per the above mentioned architecture

In [41]:

*# Define the model*

```
def cnn_model_1():  
  
    model = Sequential()  
  
    # Add layers as per the architecture mentioned above in the same sequence  
  
    model.add(Conv2D(filters=16, kernel_size=(3, 3), padding="same", input_shape=(32, 32, 1)))  
  
    model.add(LeakyReLU(0.1))  
  
    model.add(Conv2D(filters=32, kernel_size=(3, 3), padding="same"))  
  
    model.add(LeakyReLU(0.1))  
  
    model.add(MaxPool2D(pool_size=(2, 2)))  
  
    model.add(Flatten())  
  
    model.add(Dense(32))  
  
    model.add(LeakyReLU(0.1))  
  
    model.add(Dense(10, activation='softmax'))  
  
    # Compile the model  
  
    model.compile(loss='categorical_crossentropy', optimizer=tf.keras.optimizers.Adam(learning_rate = 0.001), metrics=['accuracy'])  
  
    return model
```

In [42]:

*# Build the model*

```
model_1 = cnn_model_1()
```

In [43]:

*# Print the model summary*

```
model_1.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 32, 32, 16)	160
leaky_re_lu_3 (LeakyReLU)	(None, 32, 32, 16)	0
conv2d_3 (Conv2D)	(None, 32, 32, 32)	4640
leaky_re_lu_4 (LeakyReLU)	(None, 32, 32, 32)	0
max_pooling2d_1 (MaxPooling 2D)	(None, 16, 16, 32)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 32)	262176
leaky_re_lu_5 (LeakyReLU)	(None, 32)	0
dense_3 (Dense)	(None, 10)	330
=====		
Total params: 267,306		
Trainable params: 267,306		
Non-trainable params: 0		

In [44]:

```
# Fit the model

history_model_1 = model_1.fit(
    X_train, y_train,
    epochs=20,
    validation_split=0.2,
    batch_size = 32,
    verbose=1)

Epoch 1/20
1050/1050 [=====] - 62s 58ms/step - loss: 1.1710 - accuracy: 0.6114 - val_loss: 0.6506 - val_accuracy: 0.8094
Epoch 2/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.5303 - accuracy: 0.8474 - val_loss: 0.5146 - val_accuracy: 0.8539
Epoch 3/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.4467 - accuracy: 0.8711 - val_loss: 0.4945 - val_accuracy: 0.8585
Epoch 4/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.3886 - accuracy: 0.8863 - val_loss: 0.4423 - val_accuracy: 0.8768
Epoch 5/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.3420 - accuracy: 0.8984 - val_loss: 0.4656 - val_accuracy: 0.8717
Epoch 6/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.3060 - accuracy: 0.9083 - val_loss: 0.4702 - val_accuracy: 0.8720
Epoch 7/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.2751 - accuracy: 0.9176 - val_loss: 0.4600 - val_accuracy: 0.8736
Epoch 8/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.2480 - accuracy: 0.9239 - val_loss: 0.4867 - val_accuracy: 0.8705
Epoch 9/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.2245 - accuracy: 0.9317 - val_loss: 0.4888 - val_accuracy: 0.8726
Epoch 10/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.1994 - accuracy: 0.9391 - val_loss: 0.4785 - val_accuracy: 0.8805
Epoch 11/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.1828 - accuracy: 0.9430 - val_loss: 0.5438 - val_accuracy: 0.8688
Epoch 12/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.1641 - accuracy: 0.9485 - val_loss: 0.5645 - val_accuracy: 0.8706
Epoch 13/20
1050/1050 [=====] - 62s 59ms/step - loss: 0.1476 - accuracy: 0.9546 - val_loss: 0.5580 - val_accuracy: 0.8711
Epoch 14/20
1050/1050 [=====] - 61s 59ms/step - loss: 0.1359 - accuracy: 0.9576 - val_loss: 0.5841 - val_accuracy: 0.8693
Epoch 15/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.1228 - accuracy: 0.9618 - val_loss: 0.6122 - val_accuracy: 0.8669
Epoch 16/20
1050/1050 [=====] - 61s 58ms/step - loss: 0.1092 - accuracy: 0.9654 - val_loss: 0.6472 - val_accuracy: 0.8685
Epoch 17/20
1050/1050 [=====] - 62s 59ms/step - loss: 0.0992 - accuracy: 0.9687 - val_loss: 0.6768 - val_accuracy: 0.8661
Epoch 18/20
1050/1050 [=====] - 62s 59ms/step - loss: 0.0919 - accuracy: 0.9704 - val_loss: 0.7570 - val_accuracy: 0.8631
Epoch 19/20
1050/1050 [=====] - 62s 59ms/step - loss: 0.0824 - accuracy: 0.9731 - val_loss: 0.7279 - val_accuracy: 0.8657
Epoch 20/20
1050/1050 [=====] - 62s 59ms/step - loss: 0.0829 - accuracy: 0.9731 - val_loss: 0.7561 - val_accuracy: 0.8699
```

Plotting the validation and training accuracies

Write your observations on the below plot.

In [45]:

```
# Plotting the accuracies

dict_hist = history_model_1.history

list_ep = [i for i in range(1, 21)]

plt.figure(figsize = (8, 8))

plt.plot(list_ep, dict_hist['accuracy'], ls = '--', label = 'accuracy')


plt.plot(list_ep, dict_hist['val_accuracy'], ls = '--', label = 'val_accuracy')

plt.ylabel('Accuracy')

plt.xlabel('Epochs')

plt.legend()

plt.show()
```



Observations: \_\_

- The model did poorly on the validation data; the model looks overfitting on the training data.
- The validation accuracy didn't really change after 2.5 epochs.
- The increasing rate of accuracy is lower after 2.5 epochs.

build another model and see if we can get a better model with generalized performance.

First, we need to clear the previous model's history from the Keras backend. Also, let's fix the seed again after clearing the backend.

In [46]:

```
# Clearing backend

from tensorflow.keras import backend

backend.clear_session()
```

In [47]:

```
# Fixing the seed for random number generators

np.random.seed(42)

import random

random.seed(42)

tf.random.set_seed(42)
```

## Second Model Architecture

- function that returns a sequential model with the following architecture:
  - First Convolutional layer with **16 filters and the kernel size of 3x3**. Use the **'same' padding** and provide the **input shape = (32, 32, 1)**
  - Add a **LeakyRelu layer** with the **slope equal to 0.1**
  - Second Convolutional layer with **32 filters and the kernel size of 3x3 with 'same' padding**
  - Add **LeakyRelu** with the **slope equal to 0.1**
  - Add a **max-pooling layer** with a **pool size of 2x2**
  - Add a **BatchNormalization layer**
  - Third Convolutional layer with **32 filters and the kernel size of 3x3 with 'same' padding**
  - Add a **LeakyRelu layer with the slope equal to 0.1**
  - Fourth Convolutional layer **64 filters and the kernel size of 3x3 with 'same' padding**
  - Add a **LeakyRelu layer with the slope equal to 0.1**
  - Add a **max-pooling layer** with a **pool size of 2x2**
  - Add a **BatchNormalization layer**
  - **Flatten** the output from the previous layer
  - Add a **dense layer with 32 nodes**
  - Add a **LeakyRelu layer with the slope equal to 0.1**
  - Add a **dropout layer with the rate equal to 0.5**
  - Add the final **output layer with nodes equal to the number of classes, i.e., 10** and **'softmax' as the activation function**
  - Compile the model with the **categorical\_crossentropy** loss, **adam optimizers (learning\_rate = 0.001)**, and **metric equal to 'accuracy'** . Do not fit the model here, just return the compiled model.
- Call the function `cnn_model_2` and store the model in a new variable.
- Print the summary of the model.
- Fit the model on the train data with a **validation split of 0.2, batch size = 128, verbose = 1, and epochs = 30**. Store the model building history to use later for visualization.

## Build and train the second CNN model as per the above mentioned architecture

In [52]:

```
# Define the model
```

```
def cnn_model_2():
```

```
    model = Sequential()

    # Add layers as per the architecture mentioned above in the same sequence
    model.add(Conv2D(filters=16, kernel_size=(3, 3), padding="same", input_shape=(32, 32, 1)))

    model.add(LeakyReLU(0.1))

    model.add(Conv2D(filters=32, kernel_size=(3, 3), padding="same"))

    model.add(LeakyReLU(0.1))

    model.add(MaxPool2D(pool_size=(2, 2)))

    model.add(BatchNormalization())

    model.add(Conv2D(filters=32, kernel_size=(3, 3), padding="same"))

    model.add(LeakyReLU(0.1))

    model.add(Conv2D(filters=64, kernel_size=(3, 3), padding="same"))

    model.add(LeakyReLU(0.1))

    model.add(MaxPool2D(pool_size=(2, 2)))

    model.add(BatchNormalization())

    model.add(Flatten())

    model.add(Dense(32))

    model.add(LeakyReLU(0.1))

    model.add(Dropout(0.5))

    model.add(Dense(10, activation='softmax'))

    # Compile the model

    model.compile(
        loss='categorical_crossentropy',
        optimizer=tf.keras.optimizers.Adam(learning_rate = 0.001),
        metrics=['accuracy'])
```

```
return model

# Build the model

model_2 = cnn_model_2()

# Print the summary

model_2.summary()

Model: "sequential_2"

Layer (type)                Output Shape                Param #
=====
conv2d_8 (Conv2D)           (None, 32, 32, 16)         160
leaky_re_lu_10 (LeakyReLU)  (None, 32, 32, 16)         0
conv2d_9 (Conv2D)           (None, 32, 32, 32)         4640
leaky_re_lu_11 (LeakyReLU)  (None, 32, 32, 32)         0
max_pooling2d_4 (MaxPooling (None, 16, 16, 32)    0
2D)
batch_normalization_4 (Batc (None, 16, 16, 32)         128
hNormalization)
conv2d_10 (Conv2D)          (None, 16, 16, 32)         9248
leaky_re_lu_12 (LeakyReLU)  (None, 16, 16, 32)         0
conv2d_11 (Conv2D)          (None, 16, 16, 64)         18496
leaky_re_lu_13 (LeakyReLU)  (None, 16, 16, 64)         0
max_pooling2d_5 (MaxPooling (None, 8, 8, 64)          0
2D)
batch_normalization_5 (Batc (None, 8, 8, 64)          256
hNormalization)
flatten_2 (Flatten)         (None, 4096)               0
dense_2 (Dense)             (None, 32)                 131104
leaky_re_lu_14 (LeakyReLU)  (None, 32)                 0
dropout (Dropout)           (None, 32)                 0
dense_3 (Dense)             (None, 10)                 330
=====
Total params: 164,362
Trainable params: 164,170
Non-trainable params: 192
=====

# Fit the model

history_model_2 = model_2.fit(
    X_train, y_train,
    epochs=30,
    validation_split=0.2,
    batch_size = 128,
    verbose=1)
```



```

Epoch 1/30
263/263 [=====] - 100s 379ms/step - loss: 1.3502 - accuracy: 0.5405 - val_loss: 1.7414 - val_accuracy: 0.4129
Epoch 2/30
263/263 [=====] - 100s 380ms/step - loss: 0.6604 - accuracy: 0.7981 - val_loss: 0.5678 - val_accuracy: 0.8377
Epoch 3/30
263/263 [=====] - 99s 378ms/step - loss: 0.5462 - accuracy: 0.8318 - val_loss: 0.4907 - val_accuracy: 0.8518
Epoch 4/30
263/263 [=====] - 99s 378ms/step - loss: 0.4801 - accuracy: 0.8541 - val_loss: 0.4162 - val_accuracy: 0.8855
Epoch 5/30
263/263 [=====] - 99s 378ms/step - loss: 0.4417 - accuracy: 0.8653 - val_loss: 0.4543 - val_accuracy: 0.8729
Epoch 6/30
263/263 [=====] - 99s 378ms/step - loss: 0.3988 - accuracy: 0.8798 - val_loss: 0.5975 - val_accuracy: 0.8504
Epoch 7/30
263/263 [=====] - 99s 378ms/step - loss: 0.3778 - accuracy: 0.8850 - val_loss: 0.3670 - val_accuracy: 0.8980
Epoch 8/30
263/263 [=====] - 99s 378ms/step - loss: 0.3506 - accuracy: 0.8910 - val_loss: 0.4012 - val_accuracy: 0.8889
Epoch 9/30
263/263 [=====] - 99s 378ms/step - loss: 0.3266 - accuracy: 0.8983 - val_loss: 0.3478 - val_accuracy: 0.9021
Epoch 10/30
263/263 [=====] - 100s 379ms/step - loss: 0.3006 - accuracy: 0.9057 - val_loss: 0.3743 - val_accuracy: 0.8961
Epoch 11/30
263/263 [=====] - 100s 379ms/step - loss: 0.2906 - accuracy: 0.9081 - val_loss: 0.3583 - val_accuracy: 0.9070
Epoch 12/30
263/263 [=====] - 100s 380ms/step - loss: 0.2747 - accuracy: 0.9151 - val_loss: 0.4322 - val_accuracy: 0.8960
Epoch 13/30
263/263 [=====] - 100s 381ms/step - loss: 0.2655 - accuracy: 0.9165 - val_loss: 0.3743 - val_accuracy: 0.9026
Epoch 14/30
263/263 [=====] - 100s 380ms/step - loss: 0.2503 - accuracy: 0.9206 - val_loss: 0.4257 - val_accuracy: 0.8999
Epoch 15/30
263/263 [=====] - 100s 382ms/step - loss: 0.2389 - accuracy: 0.9215 - val_loss: 0.4350 - val_accuracy: 0.9042
Epoch 16/30
263/263 [=====] - 101s 383ms/step - loss: 0.2320 - accuracy: 0.9261 - val_loss: 0.3679 - val_accuracy: 0.9105
Epoch 17/30
263/263 [=====] - 101s 383ms/step - loss: 0.2238 - accuracy: 0.9282 - val_loss: 0.4124 - val_accuracy: 0.9055
Epoch 18/30
263/263 [=====] - 100s 382ms/step - loss: 0.2167 - accuracy: 0.9297 - val_loss: 0.3827 - val_accuracy: 0.9061
Epoch 19/30
263/263 [=====] - 101s 383ms/step - loss: 0.2074 - accuracy: 0.9322 - val_loss: 0.3811 - val_accuracy: 0.9024
Epoch 20/30
263/263 [=====] - 100s 382ms/step - loss: 0.2013 - accuracy: 0.9333 - val_loss: 0.4015 - val_accuracy: 0.9067
Epoch 21/30
263/263 [=====] - 100s 382ms/step - loss: 0.1950 - accuracy: 0.9369 - val_loss: 0.4312 - val_accuracy: 0.9079
Epoch 22/30
263/263 [=====] - 101s 383ms/step - loss: 0.1877 - accuracy: 0.9394 - val_loss: 0.4362 - val_accuracy: 0.9111
Epoch 23/30
263/263 [=====] - 101s 383ms/step - loss: 0.1829 - accuracy: 0.9410 - val_loss: 0.4330 - val_accuracy: 0.9075
Epoch 24/30
263/263 [=====] - 101s 383ms/step - loss: 0.1773 - accuracy: 0.9415 - val_loss: 0.4011 - val_accuracy: 0.9118
Epoch 25/30
263/263 [=====] - 101s 384ms/step - loss: 0.1664 - accuracy: 0.9449 - val_loss: 0.4763 - val_accuracy: 0.9087
Epoch 26/30
263/263 [=====] - 101s 383ms/step - loss: 0.1670 - accuracy: 0.9445 - val_loss: 0.4698 - val_accuracy: 0.9090
Epoch 27/30
263/263 [=====] - 101s 382ms/step - loss: 0.1640 - accuracy: 0.9452 - val_loss: 0.3981 - val_accuracy: 0.9115
Epoch 28/30
263/263 [=====] - 101s 384ms/step - loss: 0.1526 - accuracy: 0.9487 - val_loss: 0.5173 - val_accuracy: 0.9032
Epoch 29/30
263/263 [=====] - 101s 383ms/step - loss: 0.1509 - accuracy: 0.9493 - val_loss: 0.4352 - val_accuracy: 0.9112
Epoch 30/30
263/263 [=====] - 101s 382ms/step - loss: 0.1483 - accuracy: 0.9506 - val_loss: 0.4499 - val_accuracy: 0.9125

```

# Plotting the validation and training accuracies

Write your observations on the below plot

In [58]:

```

# Plotting the accuracies

dict_hist = history_model_2.history

list_ep = [i for i in range(1, 31)]

plt.figure(figsize = (8, 8))

plt.plot(list_ep, dict_hist['accuracy'], ls = '--', label = 'accuracy')

plt.plot(list_ep, dict_hist['val_accuracy'], ls = '--', label = 'val_accuracy')


plt.ylabel('Accuracy')

plt.xlabel('Epochs')

```

plt.legend()

plt.show()



Observations: \_\_

- This model has reduced the overfitting as compared to the previous model but still the validation data accuracy is a bit lower than train accuracy.
- There is a rapid increase up to around 5 epochs and then seems to have very lower increase after that.
- The validation accuracy is fluctuating but generally it is also increasing with the increase in epochs.
- The test model is giving close to 90% accuracy at 30 epochs while accuracy for training model for 30 epochs, is about 95%.
- Accuracy of this model is better than the 1st model.

Predictions on the test data

- Make predictions on the test set using the second model.
- Print the obtained results using the classification report and the confusion matrix.
- Final observations on the obtained results.

Make predictions on the test data using the second model

```
In [59]:  
  
# Make prediction on the test data using model_2
```

```
test_pred = model_2.predict(X_test)  
  
test_pred = np.argmax(test_pred, axis = -1)
```

**Note:** Earlier, we noticed that each entry of the target variable is a one-hot encoded vector, but to print the classification report and confusion matrix, we must convert each entry of y\_test to a single label.

```
In [60]:  
  
# Converting each entry to single label from one-hot encoded vector  
  
y_test = np.argmax(y_test, axis = -1)
```

Write your final observations on the performance of the model on the test data.

```
In [61]:  
  
# Importing required functions  
  
from sklearn.metrics import classification_report  
  
from sklearn.metrics import confusion_matrix  
  
# Printing the classification report
```

```
print(classification_report(y_test, test_pred))
```

*# Plotting the heatmap using confusion matrix*

```
cm = confusion_matrix(y_test, test_pred)
```

```
plt.figure(figsize = (8, 5))
```

```
sns.heatmap(cm, annot = True, fmt = '.0f')
```


```
plt.ylabel('Actual')
```

```
plt.xlabel('Predicted')
```

```
plt.show()
```

	precision	recall	f1-score	support
0	0.92	0.95	0.93	1814
1	0.90	0.91	0.91	1828
2	0.94	0.91	0.93	1803
3	0.91	0.88	0.89	1719
4	0.93	0.93	0.93	1812
5	0.92	0.91	0.92	1768
6	0.91	0.90	0.90	1832
7	0.94	0.92	0.93	1808
8	0.89	0.90	0.89	1812
9	0.89	0.92	0.90	1804

accuracy		0.91	18000	
macro avg	0.91	0.91	0.91	18000
weighted avg	0.91	0.91	0.91	18000



**Final Observations: \_**

- The model gives about 90% accuracy on the test data which is comparable to the accuracy of the validation data.
- This implies that the model is giving a generalized performance.
- The recall has a very high range (82-95)% which implies that the model is good at identifying most of the objects. Model is able to identify about 95% of image 0 but can only identify only ~82% of image 3.
- Generally the model could distiguish individual digits well.
- CNN works better than ANN when identifying images.